GitHub tasks - prospective Agile-hours collection

8 real linked GitHub issues, with topic/domain/language hints. No actual worked-hour labels and no verified developer experience, so DO NOT train/test a supervised hour predictor on them. Collect verified start-of-task text and team worklogs prospectively.

Exact columns and handling

COLUMN	MEANING	POLICY
source_url, task_id	GitHub issue provenance	Trace to original
task_title, issue_summary_researcher_paraphrase	Task descriptions / researcher paraphrase	Potential future inputs; capture actual pre-start snapshot
project_domain_candidate, agile_phase_candidate_not_verified	Tentative categories	Manual verification needed
repo_language_context_not_task_verified	Repository languages	NOT necessarily task-specific
developer_experience_years	Not present	Only collect with permission, no inference
actual_hours	Empty	MUST be measured after work; never fabricated

Real row

GH-DJ-196: Django OAuth support proposal, domain Web backend/framework, repository Python context; actual_hours MISSING. Not eligible for supervised training.

Run on Windows

Open PowerShell inside the extracted project folder. Python 3.11 recommended.

```
py -3.11 -m venv .venv
.\.venv\Scripts\python.exe -m pip install -r requirements-base.txt
.\.venv\Scripts\python.exe audit_github.py
```

What is and is not validated

The command audits rows and saves results/ (CSV/JSON). There is no 70/15/15 split: insufficient compatible text labels. No fine-tuning or supervised testing is scientifically possible with this file alone.

For new records, collect issue title/description/version BEFORE work; optional project domain, task phase, repository languages and consented experience; then after work sum verified person-worklog minutes / 60. Track contributor counts and corrections. Do not mistake wall-clock issue age or commit counts for hours.